


The learned LoRA baseline exceeds 90% on both fact groups

Each row compares all prompts with held-out paraphrases; the learned starting model is highlighted.

● All prompts ● Held-out paraphrases Highlighted learned LoRA model

Forget-fact accuracy

Retain-fact accuracy



Before unlearning, high forget accuracy confirms learning. During unlearning, lower forget accuracy becomes better.